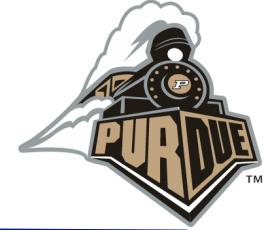


Toward Generating Reducible Replay Logs

Kyu Hyung Lee, Yunhui Zheng,
Nick Sumner, Xiangyu Zhang

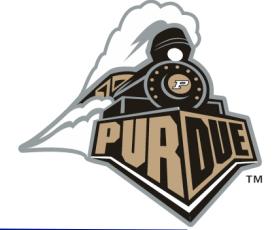
PLDI 2011, June 4 - 7





Introduction

- Recording and replay technique
 - Low overhead on logging phase
 - Fast and deterministic failure replay
- Fast replay techniques for long running applications
 - Checkpointing
 - Frequent checkpoints : high recording overhead
 - Infrequent checkpoints : long running replay
 - Log Reduction



Introduction

- Log reduction
 - Reduces a replay execution
 - Remove events from the replay log, but can still reproduce the failure
 - Limitations
 - Very expensive analysis to detect dependences between events
 - Does not guarantee replayability



Log Reduction – conceptual example

```
1 int count=0;  
2 Node *list = null;  
3 while(read(&name, &age)) {  
4     if(age > 18) {  
5         n = new Node(name);  
6         n->next = list;  
7         list = n;  
8         print(name);  
9     }  
10    count++;  
11}  
12 print("count="+count);  
13 Assert(list->next==null);
```

**Read personal information
(name & age) and store the adult
information to a linked-list**



Log Reduction – conceptual example

*/*Read personal information (name & age) and store the adult information to a linked-list*/*

```
1 int count=0;  
2 Node *list = null;  
3 while(read(&name, &age)) {  
4     if(age > 18) {  
5         n = new Node(name);  
6         n->next = list;  
7         list = n;  
8         print(name);  
9     }  
10    count++;  
11}  
10 print("count="+count);  
11Assert(list->next==null);
```

Assert the linked-list has only one element



Log Reduction – conceptual example

*/*Read personal information (name & age) and store the adult information to a linked-list*/*

```
1 int count=0;
2 Node *list = null;
3 while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11}
12 print("count="+count);
13 Assert(list->next==null);
/* Assert the list has one element */
```

Execute this program on top of the logging tool to record non-deterministic system calls



Log Reduction – conceptual example

*/*Read personal information (name & age) and store the adult information to a linked-list*/*

```
1 int count=0;
2 Node *list = null;
3 while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11}
10 print("count="+count);
11 Assert(list->next==null);
/* Assert the list has one element */
```

Execute this program on top of the logging tool to record non-deterministic system calls



Log Reduction – conceptual example

```
1 int count=0;
2 Node *list = null;
3 while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11}
12 print("count="+count);
13 Assert(list->next==null);
/* Assert the list has one element */
```

list
↓
NULL



Log Reduction – conceptual example

```
1 int count=0;
2 Node *list = null;
3 while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11}
10 print("count="+count);
11 Assert(list->next==null);
/* Assert the list has one element */
```

list
↓
NULL

LOG
3 Read(Brian, 25)



Log Reduction – conceptual example

```
1 int count=0;
2 Node *list = null;
3 while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11}
12 print("count="+count);
13 Assert(list->next==null);
/* Assert the list has one element */
```



LOG

3 Read(Brian, 25)



Log Reduction – conceptual example

```
1 int count=0;
2 Node *list = null;
3 while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11}
12 print("count="+count);
13 Assert(list->next==null);
/* Assert the list has one element */
```



LOG

- 3 Read(Brian, 25)
- 8 Write(Brian)



Log Reduction – conceptual example

```
1 int count=0;
2 Node *list = null;
3 while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11
12    print("count=" + count);
13    Assert(list->next==null);
14    /* Assert the list has one element */
```



LOG

- 3 Read(Brian, 25)
- 8 Write(Brian)



Log Reduction – conceptual example

```
1 int count=0;
2 Node *list = null;
3 while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11}
12 print("count="+count);
13 Assert(list->next==null);
/* Assert the list has one element */
```



LOG

- 3 Read(Brian, 25)
- 8 Write(Brian)
- 3 Read(John, 15)



Log Reduction – conceptual example

```
1 int count=0;
2 Node *list = null;
3 while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11}
12 print("count="+count);
13 Assert(list->next==null);
/* Assert the list has one element */
```



LOG

- 3 Read(Brian, 25)
- 8 Write(Brian)
- 3 Read(John, 15)



Log Reduction – conceptual example

```
1 int count=0;
2 Node *list = null;
3 while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11
12    print("count=" + count);
13    Assert(list->next==null);
14    /* Assert the list has one element */
```



LOG

- 3 Read(Brian, 25)
- 8 Write(Brian)
- 3 Read(John, 15)



Log Reduction – conceptual example

```
1 int count=0;
2 Node *list = null;
3 while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11}
12 print("count="+count);
13 Assert(list->next==null);
/* Assert the list has one element */
```



LOG

3	Read(Brian, 25)
8	Write(Brian)
3	Read(John, 15)
3	Read(Tom, 21)
8	Write(Tom)
3	Read(Betty, 30)
8	Write(Betty)
3	Read(Amy, 6)
3	Read(Alex, 15)



Log Reduction – conceptual example

```
1 int count=0;
2 Node *list = null;
3 while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11}
10 print("count="+count);
11 Assert(list->next==null);
/* Assert the list has one element */
```



LOG

```
3 Read(Brian, 25)
8 Write(Brian)
3 Read(John, 15)
3 Read(Tom, 21)
8 Write(Tom)
3 Read(Betty, 30)
8 Write(Betty)
3 Read(Amy, 6)
3 Read(Alex,15)
3 Read(EOF)
```



Log Reduction – conceptual example

```
1 int count=0;
2 Node *list = null;
3 while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11}
10 print("count=" + count);
11 Assert(list->next==null);
/* Assert the list has one element */
```



LOG

3	Read(Brian, 25)
8	Write(Brian)
3	Read(John, 15)
3	Read(Tom, 21)
8	Write(Tom)
3	Read(Betty, 30)
8	Write(Betty)
3	Read(Amy, 6)
3	Read(Alex, 15)
3	Read(EOF)
10	Write(count=6)



Log Reduction – conceptual example

```
1 int count=0;
2 Node *list = null;
3 while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11}
10 print("count="+count);
11Assert(list->next==null);
/* Assert the list has one element */
```



LOG

```
3 Read(Brian, 25)
8 Write(Brian)
3 Read(John, 15)
3 Read(Tom, 21)
8 Write(Tom)
3 Read(Betty, 30)
8 Write(Betty)
3 Read(Amy, 6)
3 Read(Alex,15)
3 Read(EOF)
10 Write(count=6)
11 Assertion Fails
```



Log Reduction – conceptual example

```
1 int count=0;  
2 Node *list = null;  
3 while(read(&name, &age)) {  
4     if(age > 18) {  
5         n = new Node(name);  
6         n->next = list;  
7         list = n;  
8         print(name);  
9     }  
10    count++;  
11    /* Assert the list has one element */  
11Assert(list->next==null);
```



LOG

3	Read(Brian, 25)
8	Write(Brian)
3	Read(John, 15)
3	Read(Tom, 21)
8	Write(Tom)
3	Read(Betty, 30)
8	Write(Betty)
3	Read(Amy, 6)
3	Read(Alex, 15)
3	Read(EOF)
10	Write(count=6)

Replay
Criterion

11 Assertion Fails



Log Reduction – conceptual example

```
1 int count=0;
2 Node *list = null;
3 while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11}
10 print("count="+count);
11 Assert(list->next==null);
/* Assert the list has one element */
```



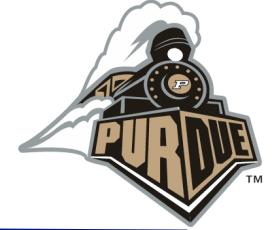
Replay
Criterion

LOG

```
3 Read(Brian, 25)
8 Write(Brian)
3 Read(John, 15)
3 Read(Tom, 21)
8 Write(Tom)
3 Read(Betty, 30)
8 Write(Betty)
3 Read(Amy, 6)
3 Read(Alex, 15)
3 Read(EOF)
10 Write(count=6)
```

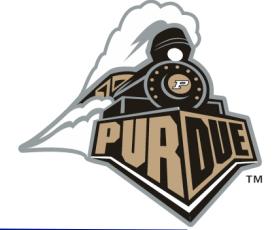
11 Assertion Fails

Expensive analysis to find reduction at
the first replay attempt



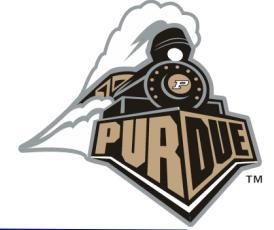
Goal

- Compiler based technique that generates a reducible replay log
 - Instrument programs to collect minimal additional information
- Reduction achieved though analyzing just the log
- Guarantee replayability of the reduced log



Challenges

- Reduction may induce a different control flow path

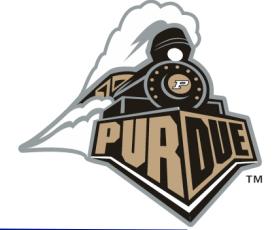


Challenges

- Reduction may induce a different control flow path

```
1 while(read(&x))  
2   if(x) write(...);
```

Read(0)
Read(10)
Write(...)

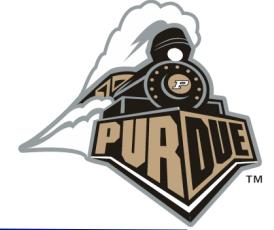


Challenges

- Reduction may induce a different control flow path

```
1 while(read(&x))  
2   if(x) write(...);
```

Read(0)
~~Read(10)~~
Write(...) // Replay FAIL!



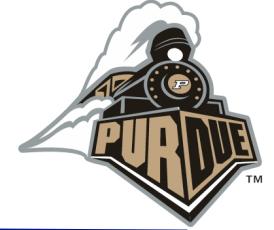
Challenges

- Reduction may induce a different control flow path

```
1 while(read(&x))  
2   if(x) write(...);
```

Read(0)
~~Read(10)~~
Write(...) // Replay FAIL!

Reduced log can not properly aligned with the execution



Challenges

- Reduction may induce a different control flow path
- Reduction may change variable values, leading to inconsistency



Challenges

- Reduction may induce a different control flow path
- Reduction may change variable values, leading to inconsistency

```
1 While(read(&size))  
2   socket_read(string, size);
```

```
Read(5)  
Socket_read("abcde")  
Read(3)  
Socket_read("fgh")
```

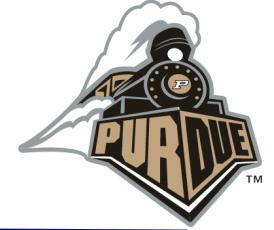


Challenges

- Reduction may induce a different control flow path
- Reduction may change variable values, leading to inconsistency

```
1 While(read(&size))  
2   socket_read(string, size);
```

Read(5)
~~Socket_read("abcde")~~
~~Read(3)~~
Socket_read("fgh")



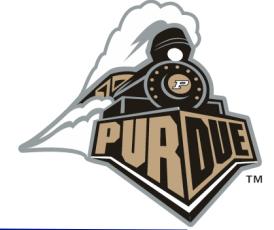
Challenges

- Reduction may induce a different control flow path
- Reduction may change variable values, leading to inconsistency

```
1 While(read(&size))  
2   socket_read(string, size);
```

Read(5)
~~Socket_read("abcde")~~
~~Read(3)~~
Socket_read("fgh") // Replay FAIL!

Number of bytes are changed
that are supposed to be read



Challenges

- Reduction may induce a different control flow path
- Reduction may change variable values, leading to inconsistency
- Structural constraints



Challenges

- Reduction may induce a different control flow path
- Reduction may change variable values, leading to inconsistency
- Structural constraints

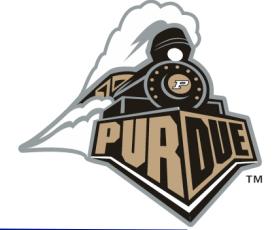
1	read(&x);	1	R(10)
2	print("x=" + x);	2	W(x=10)
3	read(&y);	3	R(20)
4	print("y=" + y);	4	W(y=20)
5	read(&z);	5	R(30)
6	print("z=" + z);	6	W(z=30)



Challenges

- Reduction may induce a different control flow path
- Reduction may change variable values, leading to inconsistency
- Structural constraints

1	read(&x);	1	R(10)
2	print("x=" + x);	2	W(x=10)
3	read(&y);	3	R(20)
4	print("y=" + y);	4	W(y=20)
5	read(&z);	5	R(30)
6	print("z=" + z);	6	W(z=30)



Challenges

- Reduction may induce a different control flow path
- Reduction may change variable values, leading to inconsistency
- Structural constraints

1	read(&x);	1	R(10)
2	print("x=" + x);	2	W(x=10)
3	read(&y);	3	R(20)
4	print("y=" + y);	4	W(y=20)
5	read(&z);	5	R(30)
6	print("z=" + z);	6	W(z=30)

Cannot reduce any event due to the straight line structure



Our Solutions

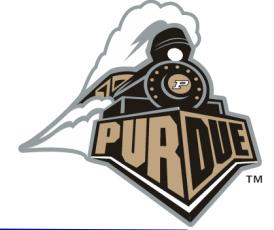
Divide execution into Units
“User annotation”

Instrumentation and optimization
“Static analysis”

Execute instrumented program

Reduction
“analyze just the log”

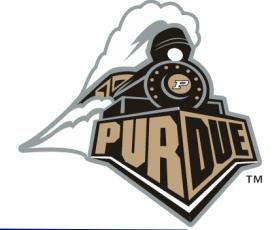
Replay reduced log



Execution UNIT

- User annotates “[UNIT]” on the event processing loop

```
1 int count=0;
2 Node *list = null;
3 while(read(&name,
4   &age)) {
5     if(age > 18) {
6         n = new Node(name);
7         n->next = list;
8         list = n;
9         print(name);
10    }
11    count++;
12}
13 print("count="+count);
14 Assert(list->next==null);
```



Execution UNIT

- User annotates “[UNIT]” on the event processing loop

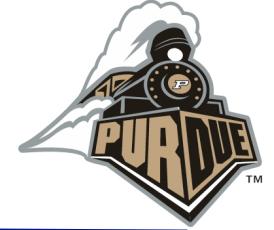
```
1 int count=0;
2 Node *list = null;
3 while(read(&name,
4   &age)) {
5     if(age > 18) {
6         n = new Node(name);
7         n->next = list;
8         list = n;
9         print(name);
10    }
11    count++;
12}
13 print("count="+count);
14 Assert(list->next==null);
```



Execution UNIT

- User annotates “[UNIT]” on the event processing loop

```
1 int count=0;
2 Node *list = null;
3 [UNIT] while(read
4 (&name, &age)) {
5     if(age > 18) {
6         n = new Node(name);
7         n->next = list;
8         list = n;
9         print(name);
10    }
11    count++;
12}
13 print("count="+count);
14 Assert(list->next==null);
```



Execution UNIT

- User annotates “[UNIT]” on the event processing loop
- Unit is an iteration of an event processing loop
- Reduction is only carried out at the unit level

```
1 int count=0;
2 Node *list = null;
3 [UNIT] while(read
4 (&name, &age)) {
5   if(age > 18) {
6     n = new Node(name);
7     n->next = list;
8     list = n;
9     print(name);
10  }
11  count++;
12
13 print("count="+count);
14 Assert(list->next==null);
```



Execution UNIT

■ Why Unit?

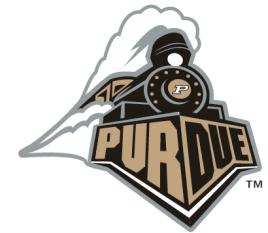
- Fewer dependences
- Naturally solve the problem of structural constraint
- Most long running programs are dominated by the event processing loop

```
1 int count=0;
2 Node *list = null;
3 [UNIT] while(read
4 (&name, &age)) {
5   if(age > 18) {
6     n = new Node(name);
7     n->next = list;
8     list = n;
9     print(name);
10  }
11  count++;
12
13  print("count=" + count);
14  Assert(list->next==null);
```



Instrumentation

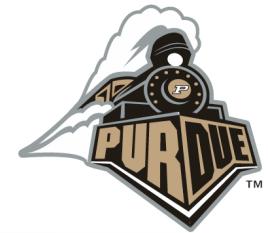
```
1 int count=0;
2 Node *list = null;
3 [UNIT] while(read(&name,
4   &age)) {
5   if(age > 18) {
6     n = new Node(name);
7     n->next = list;
8     list = n;
9     print(name);
10  }
11  count++;
12 }
13 print("count="+count);
14 Assert(list->next==null);
```



Instrumentation

```
1 int count=0;
2 Node *list = null;
3 [UNIT] while(read(&name,
4     &age)) {
5     if(age > 18) {
6         n = new Node(name);
7         n->next = list;
8         list = n;
9         print(name);
10    }
11    count++;
12 }
13 print("count="+count);
14 Assert(list->next==null);
```

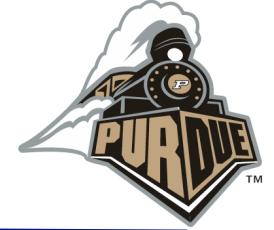
- Find all variables which possibly have inter UNIT dependence
 - Static analysis by LLVM



Instrumentation

```
1 int count=0;
2 Node *list = null;
3 [UNIT]while(read(&name,
4     &age)) {
5     if(age > 18) {
6         n = new Node(name);
7         n->next = list;
8         list = n;
9         print(name);
10    }
11    count++;
12}
13 print("count="+count);
14 Assert(list->next==null);
```

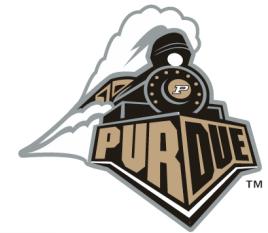
- Find all variables which possibly have inter UNIT dependence
 - Static analysis by LLVM



Instrumentation

```
1 int count=0;
2 Node *list = null;
3 [UNIT] while(read(&name,
4     &age) {
5     if(age > 18) {
6         n = new Node(name);
7         n->next = list;
8         list = n;
9         print(name);
10    }
11    count++;
12 }
13 print("count="+count);
14 Assert(list->next==null);
```

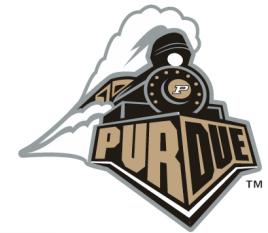
- Find all variables which possibly have inter UNIT dependence
 - Static analysis by LLVM



Instrumentation

```
1 int count=0;
2 Node *list = null;
3 [UNIT]while(read(&name,
&age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11}
12 print("count="+count);
13 Assert(list->next==null);
```

- Instrument the program
 - Log accesses of inter-unit shared variable



Basic Reduction

```
1 int count=0;
2 Node *list = null;
3 [UNIT] while(read(&name,
&age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
9     count++;
10    print("count="+count);
11    Assert(list->next==null);
```

Original log

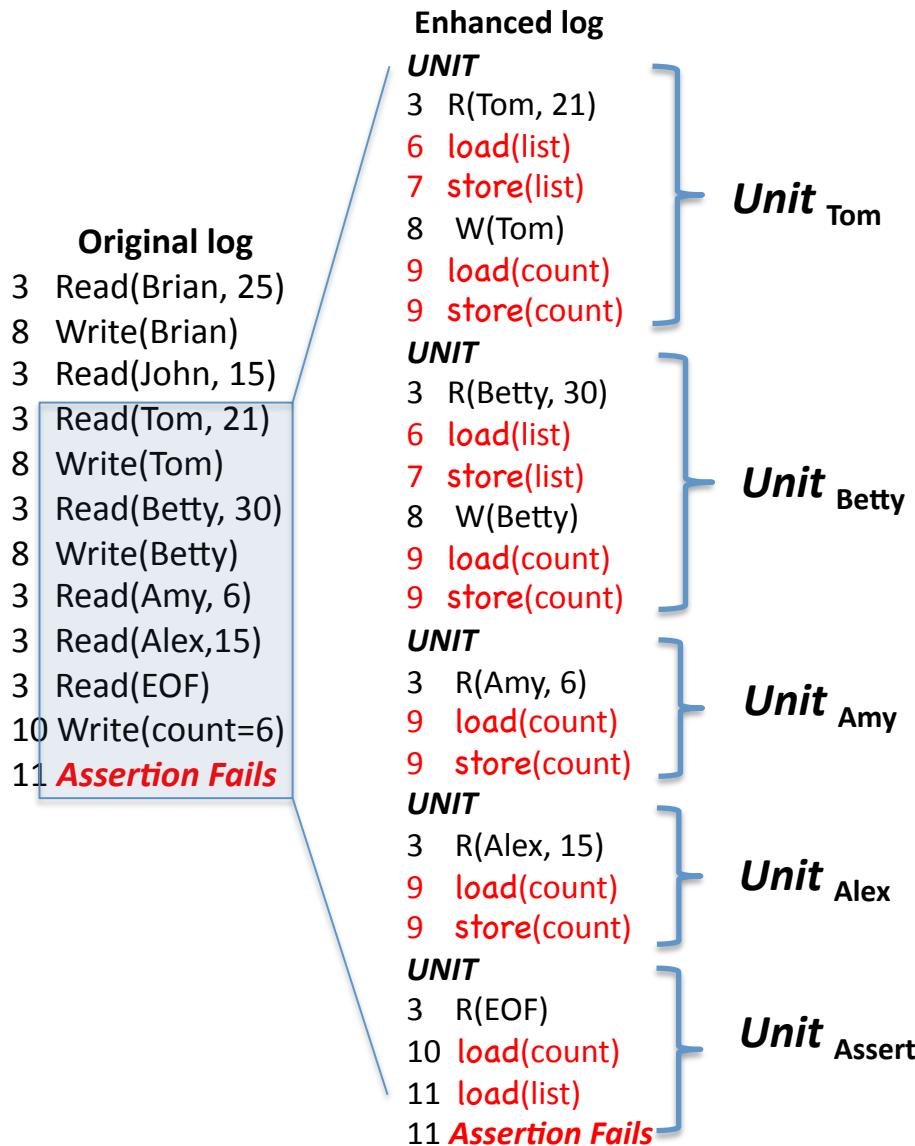
```
3 Read(Brian, 25)
8 Write(Brian)
3 Read(John, 15)
3 Read(Tom, 21)
8 Write(Tom)
3 Read(Betty, 30)
8 Write(Betty)
3 Read(Amy, 6)
3 Read(Alex,15)
3 Read(EOF)
10 Write(count=6)
11 Assertion Fails
```

Unit Tom

```
3 R(Tom, 21)
6 load(list)
7 store(list)
8 W(Tom)
9 load(count)
9 store(count)
```

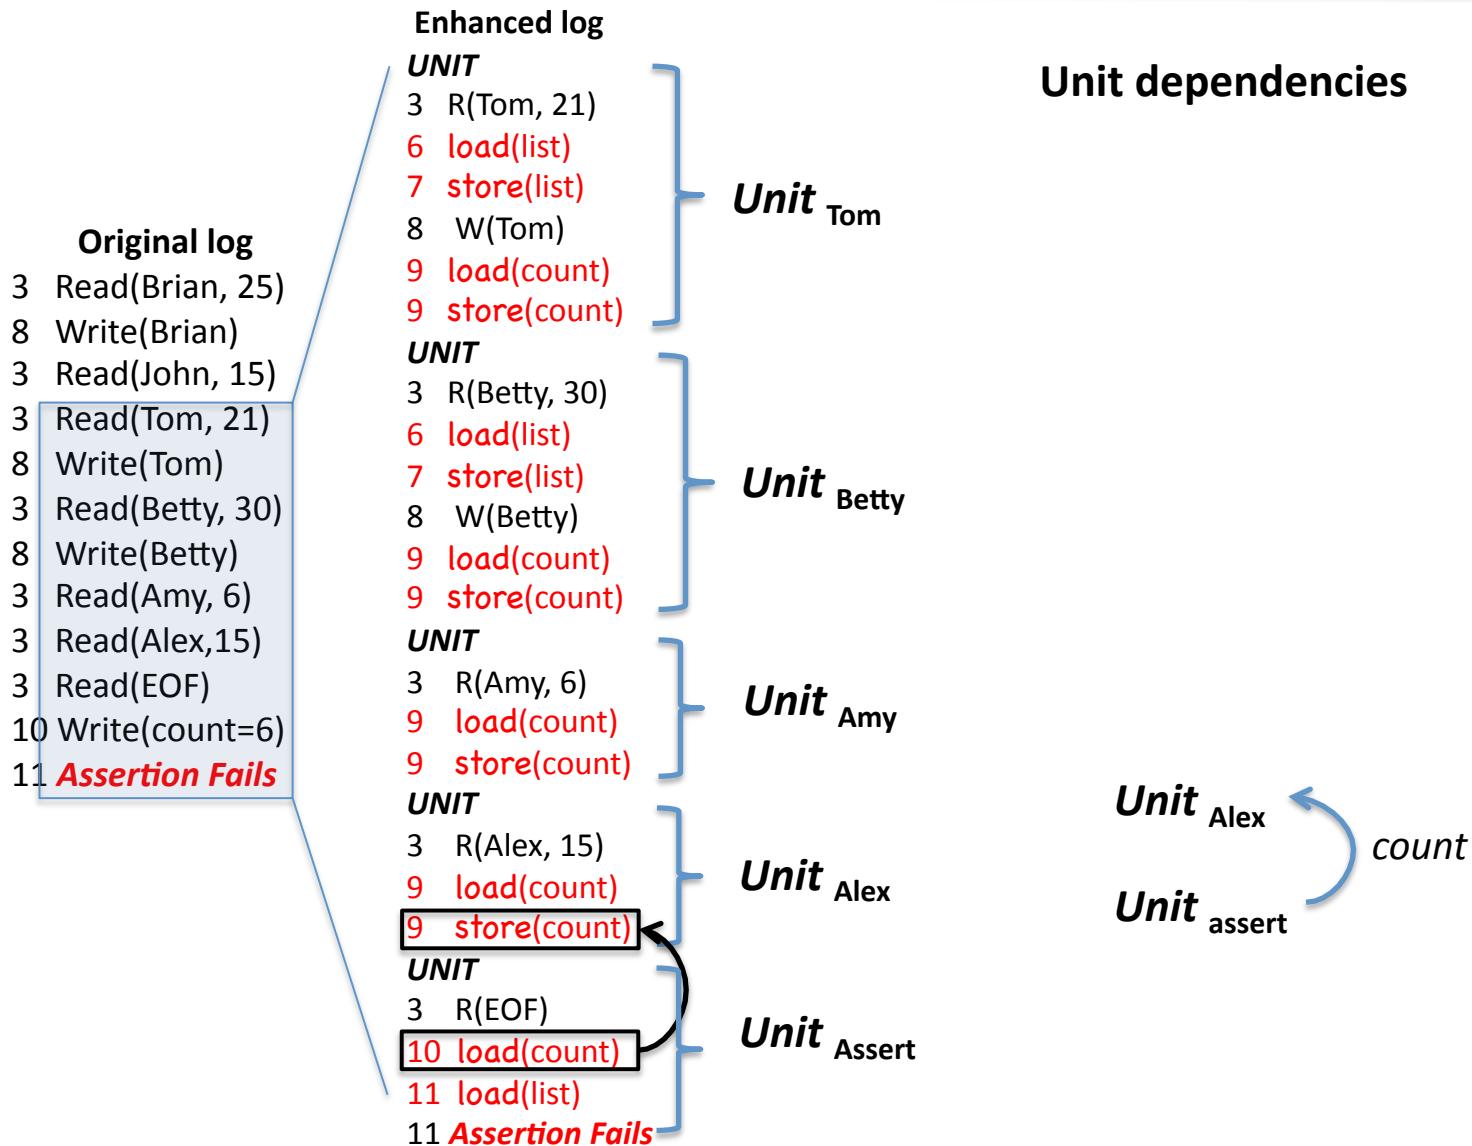


Basic Reduction



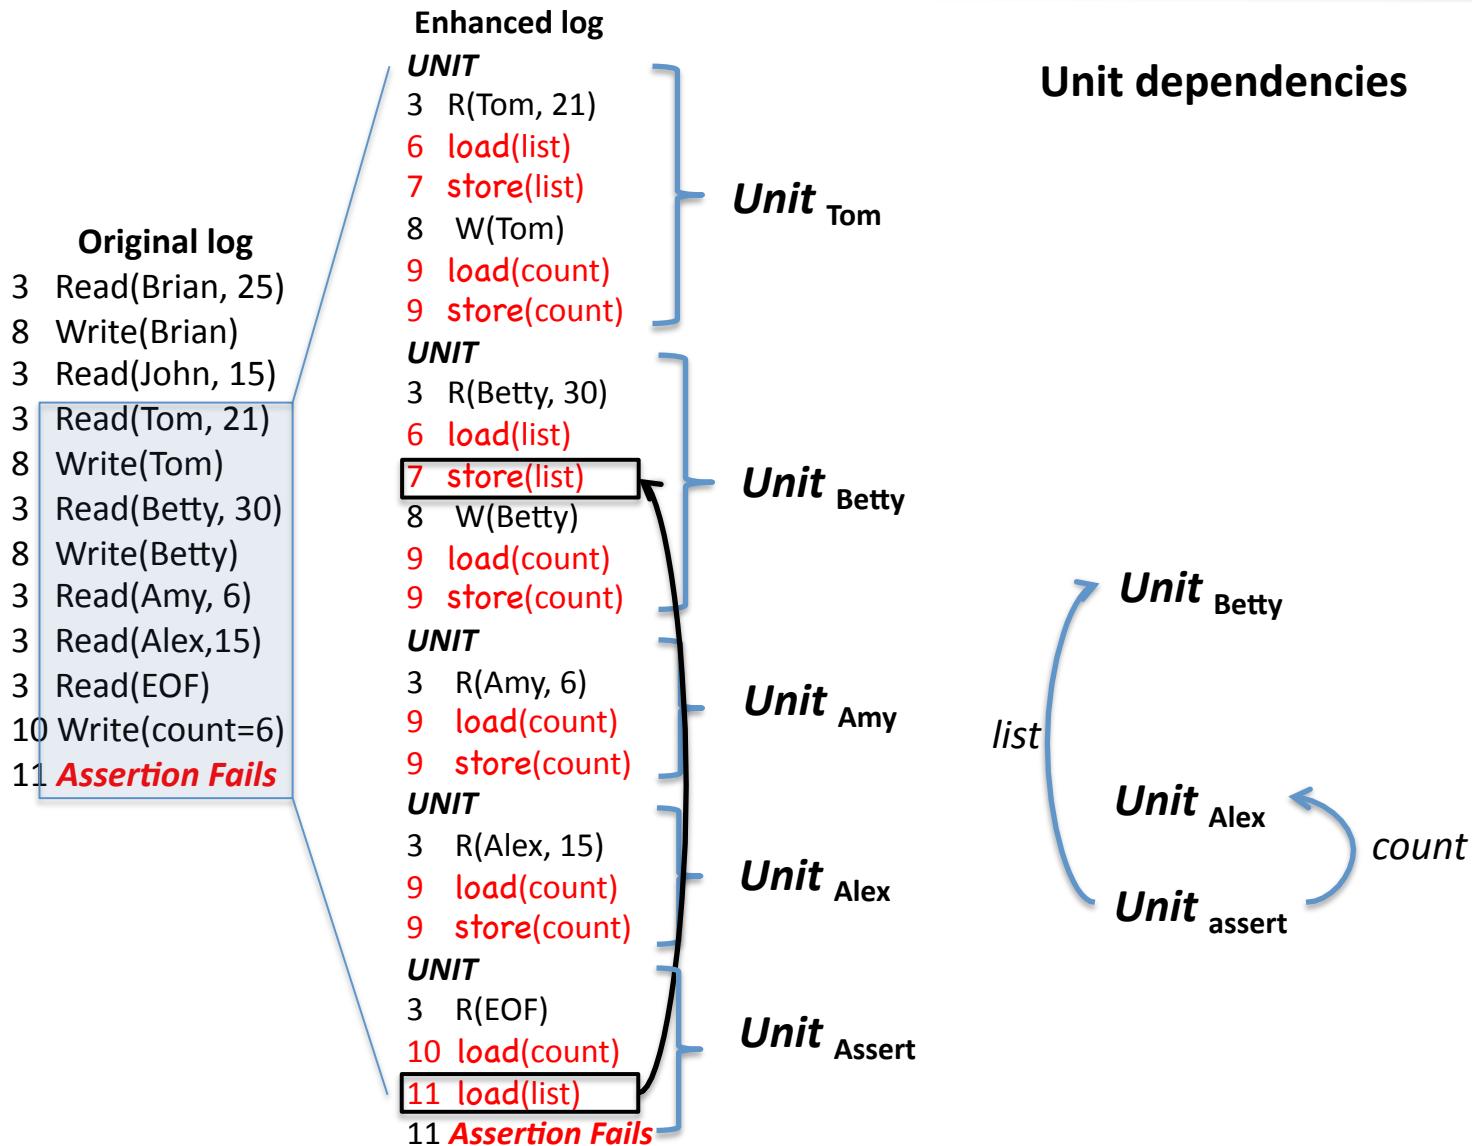


Basic Reduction



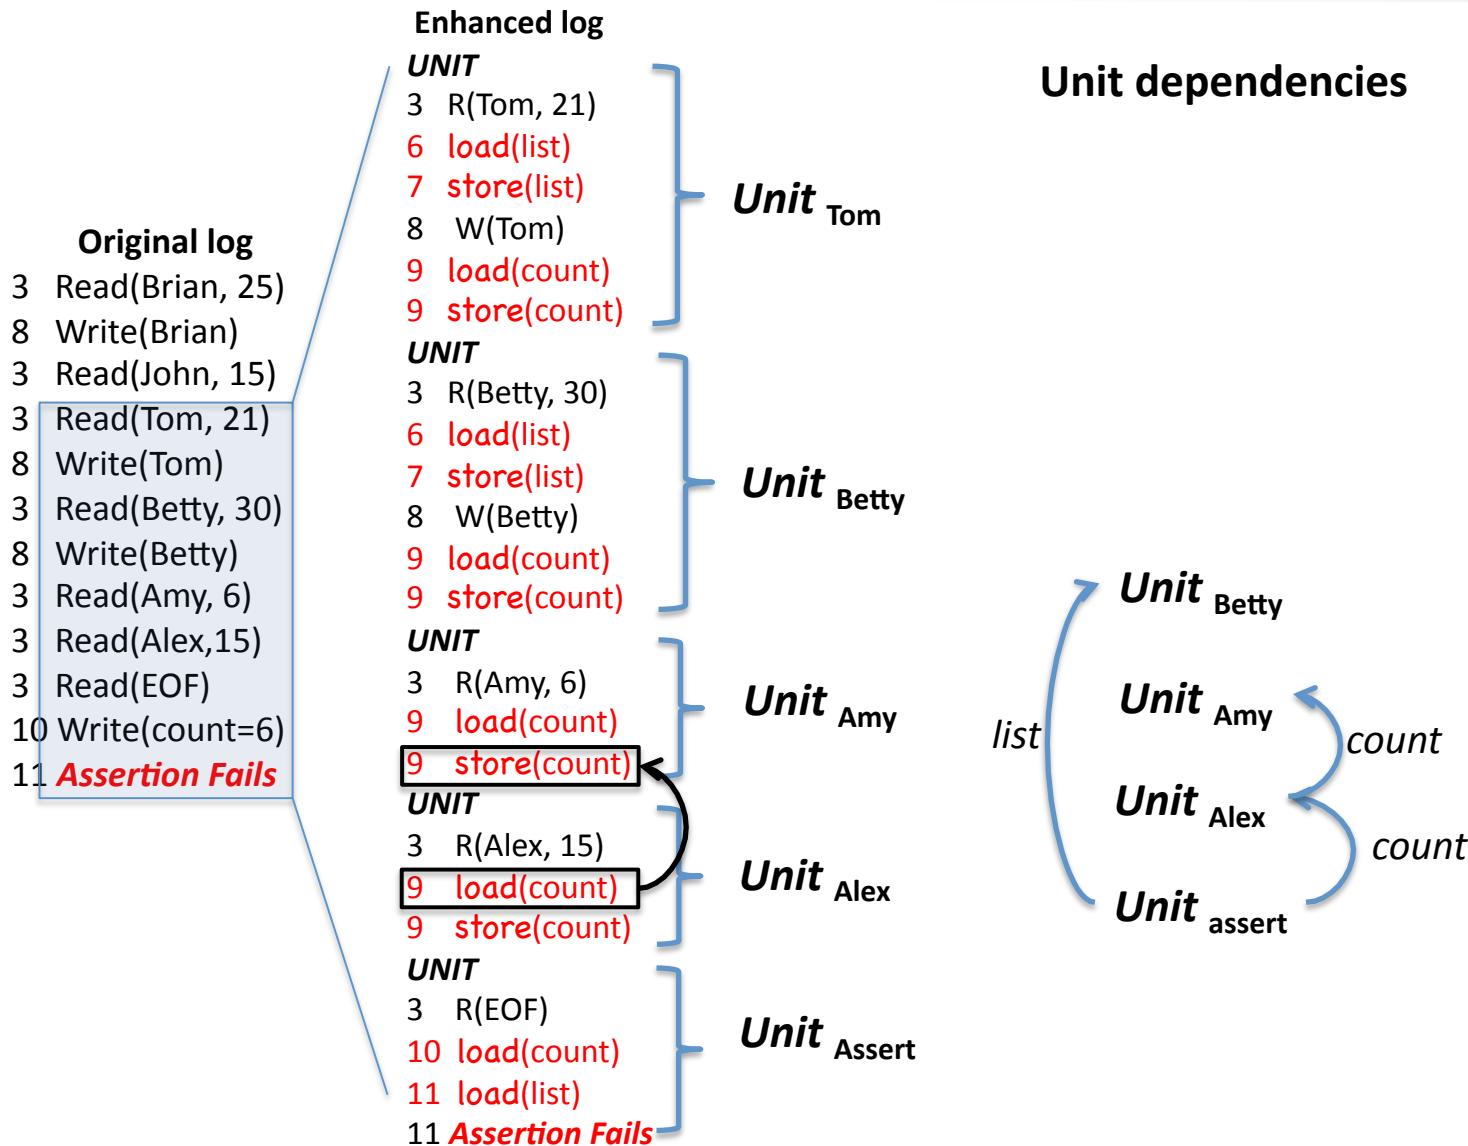


Basic Reduction



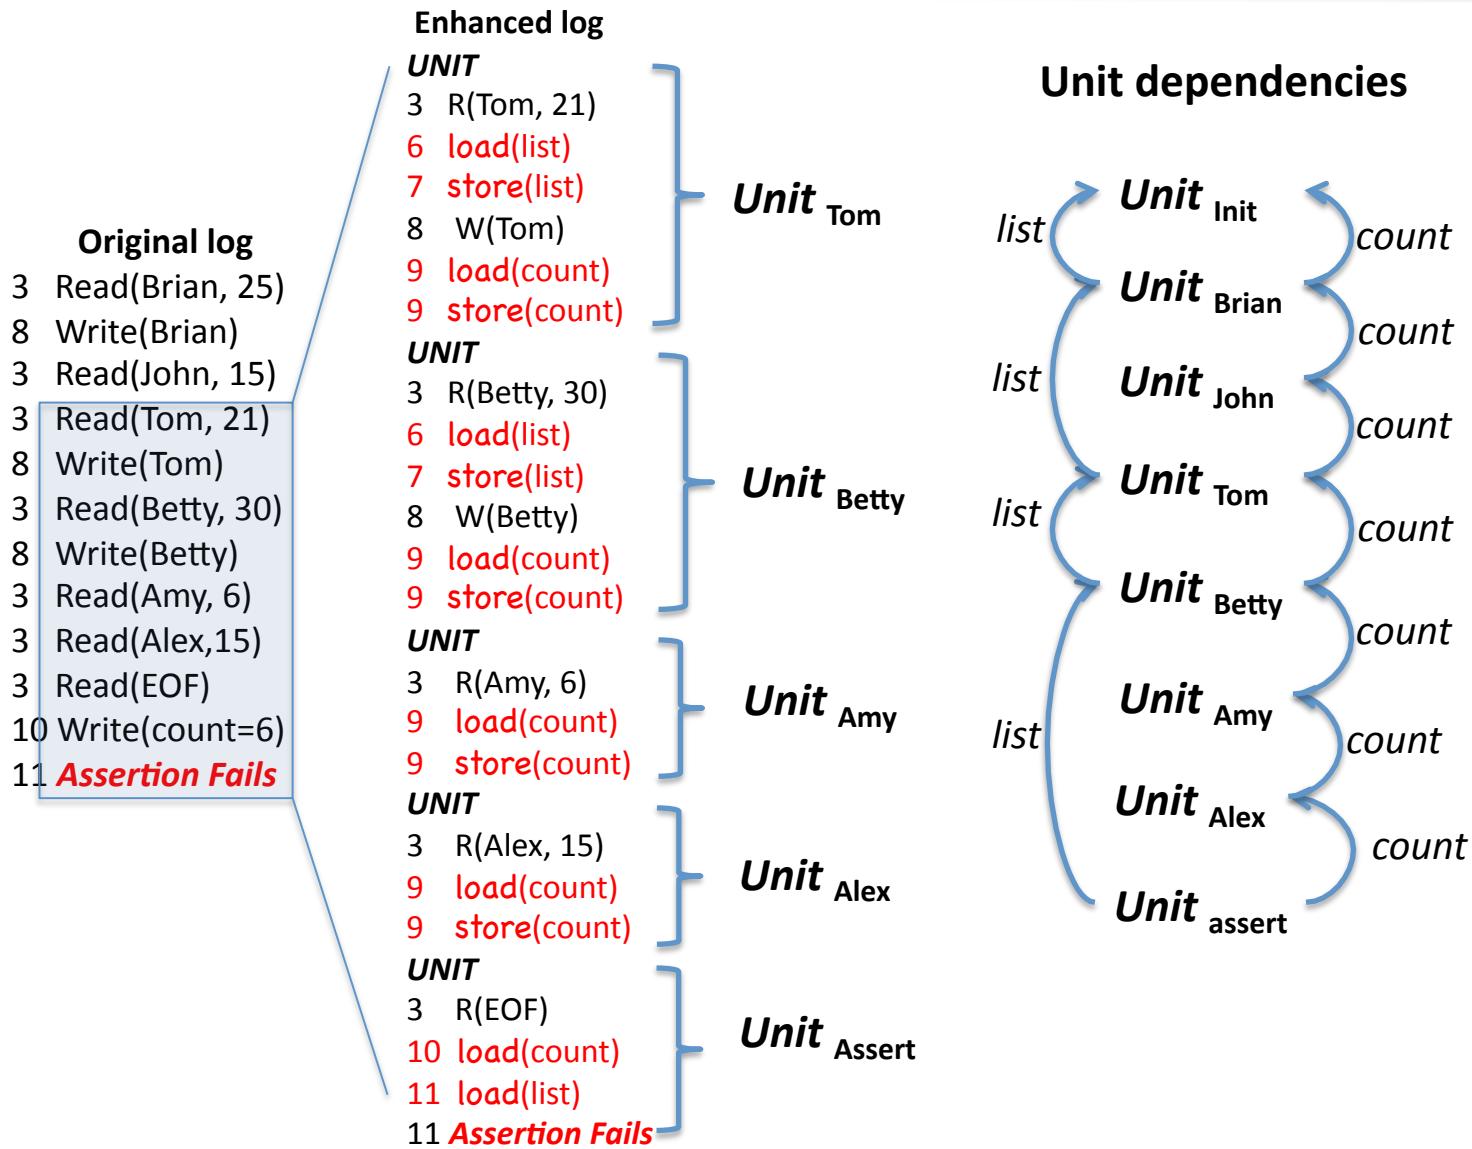


Basic Reduction



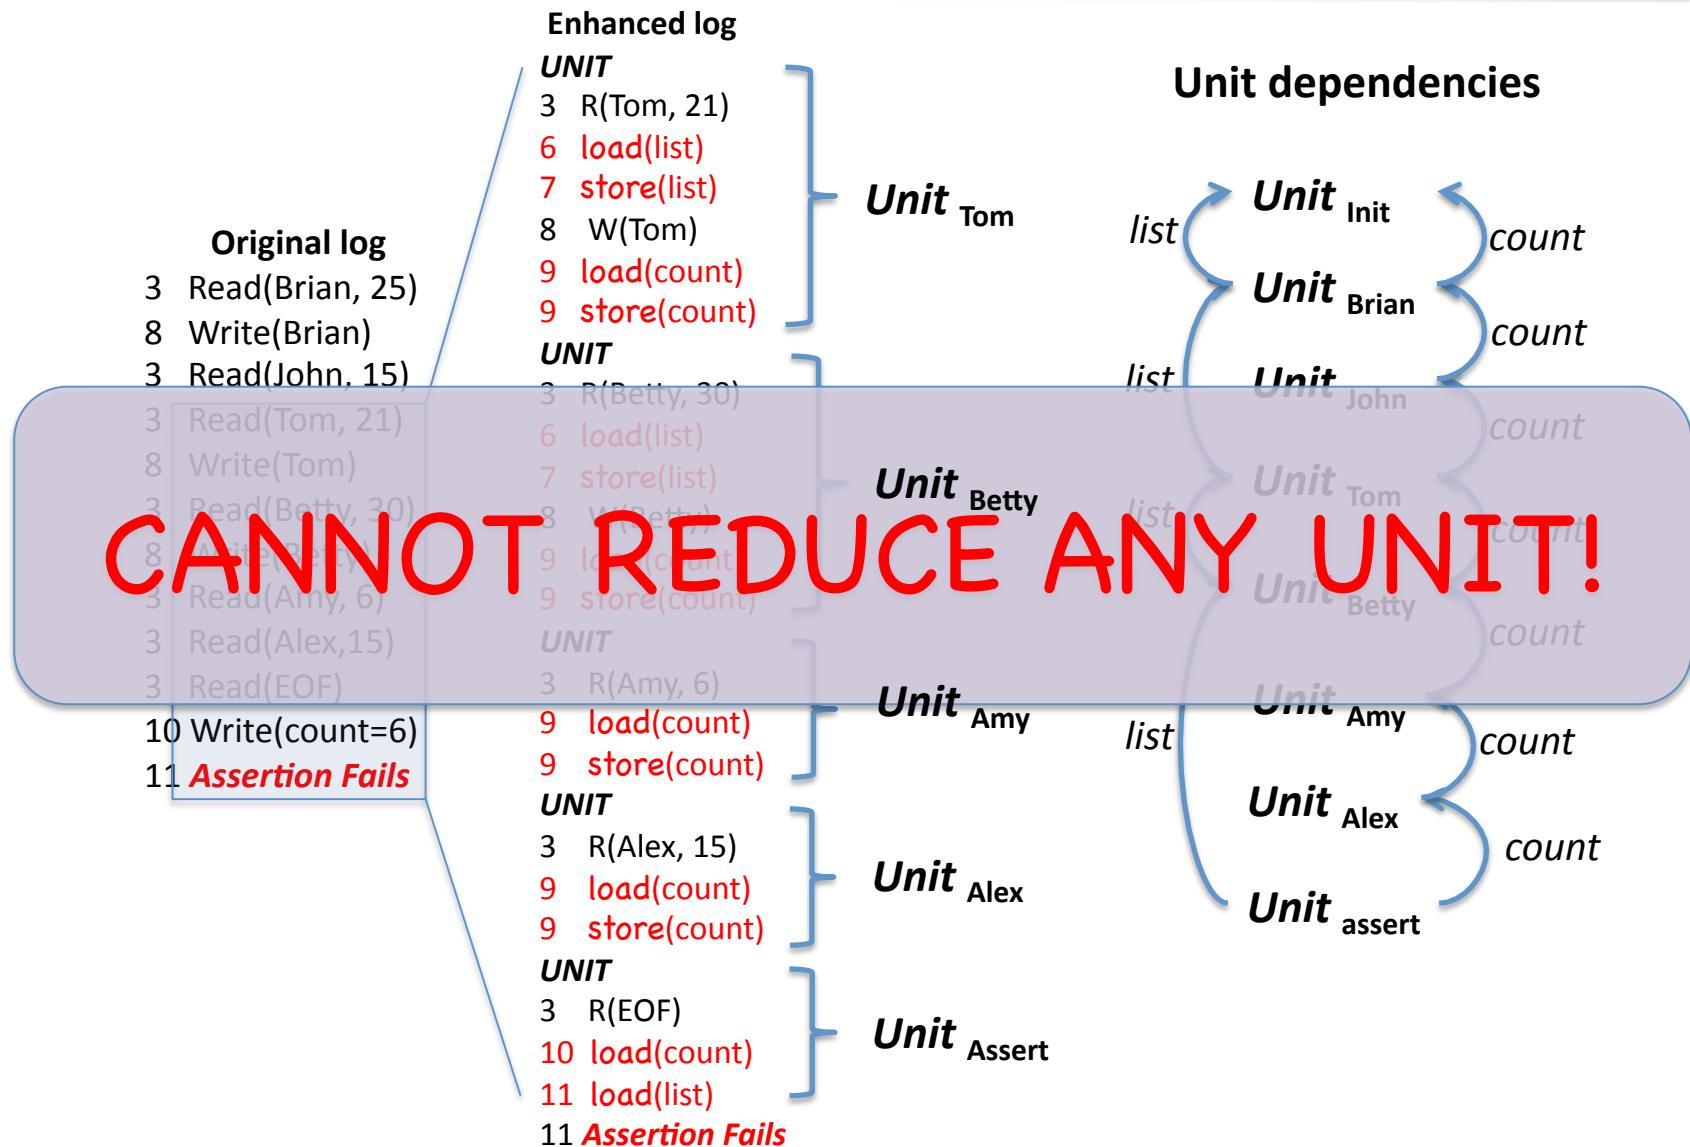


Basic Reduction





Basic Reduction



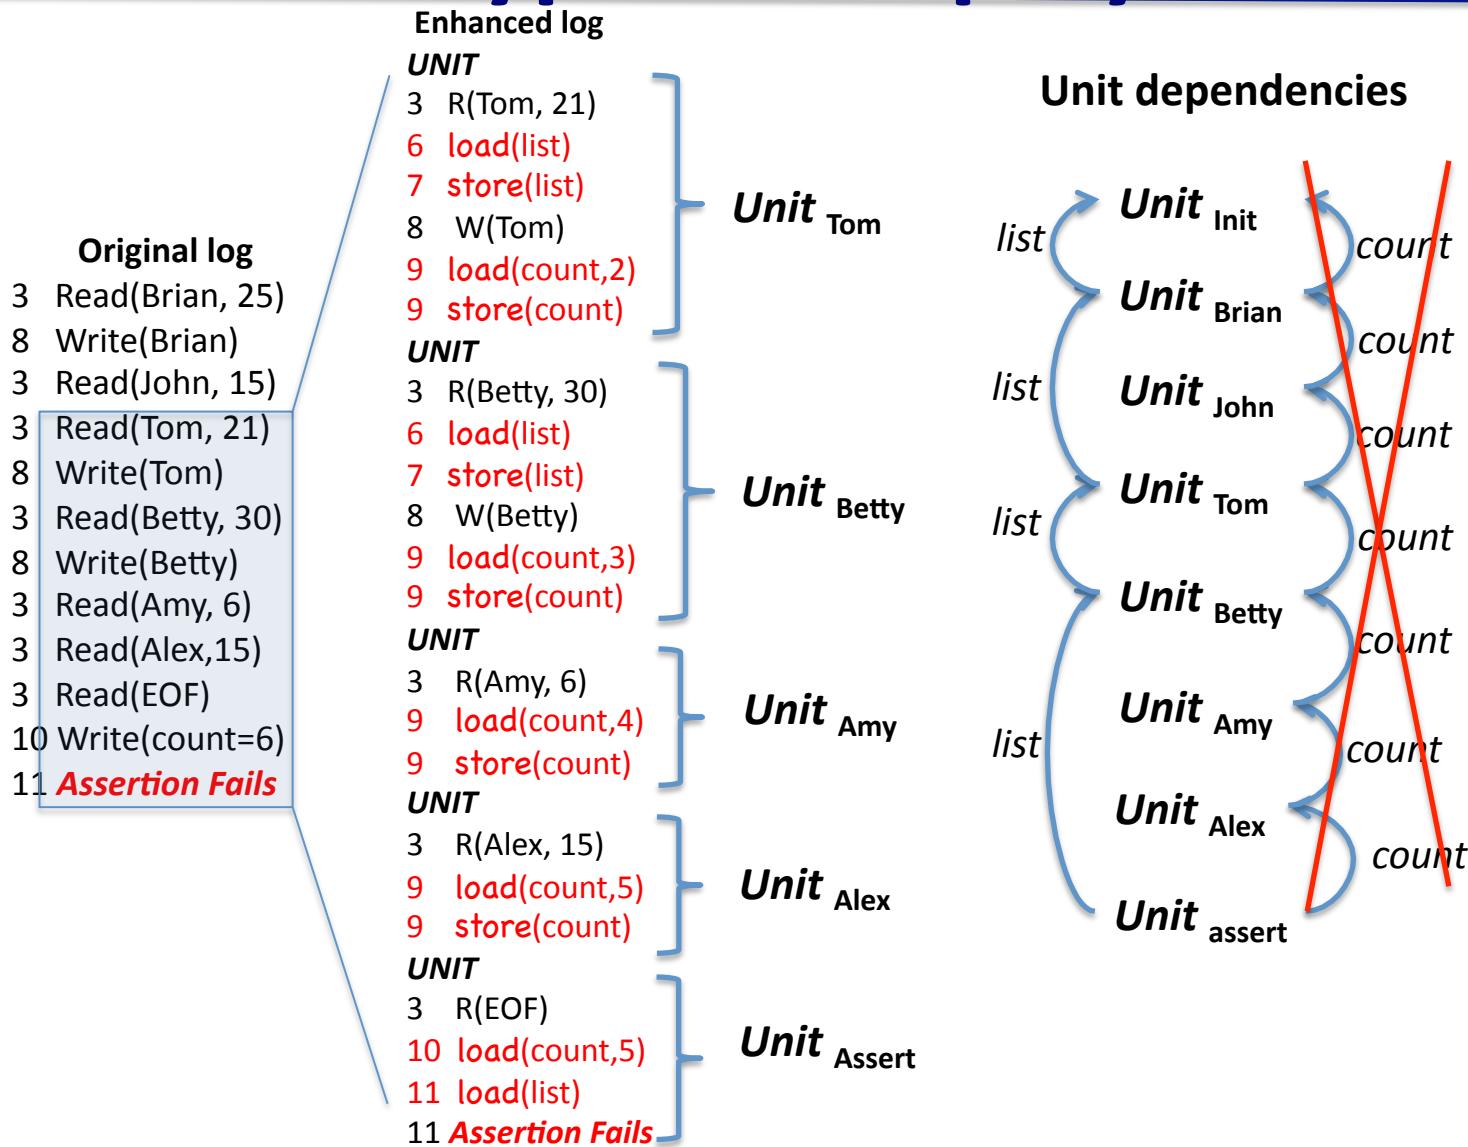


Weave two types of replay

- Replay by value
 - For primitive types
 - Record accesses and value
- Replay by dependence
 - For pointer types
 - Record accesses

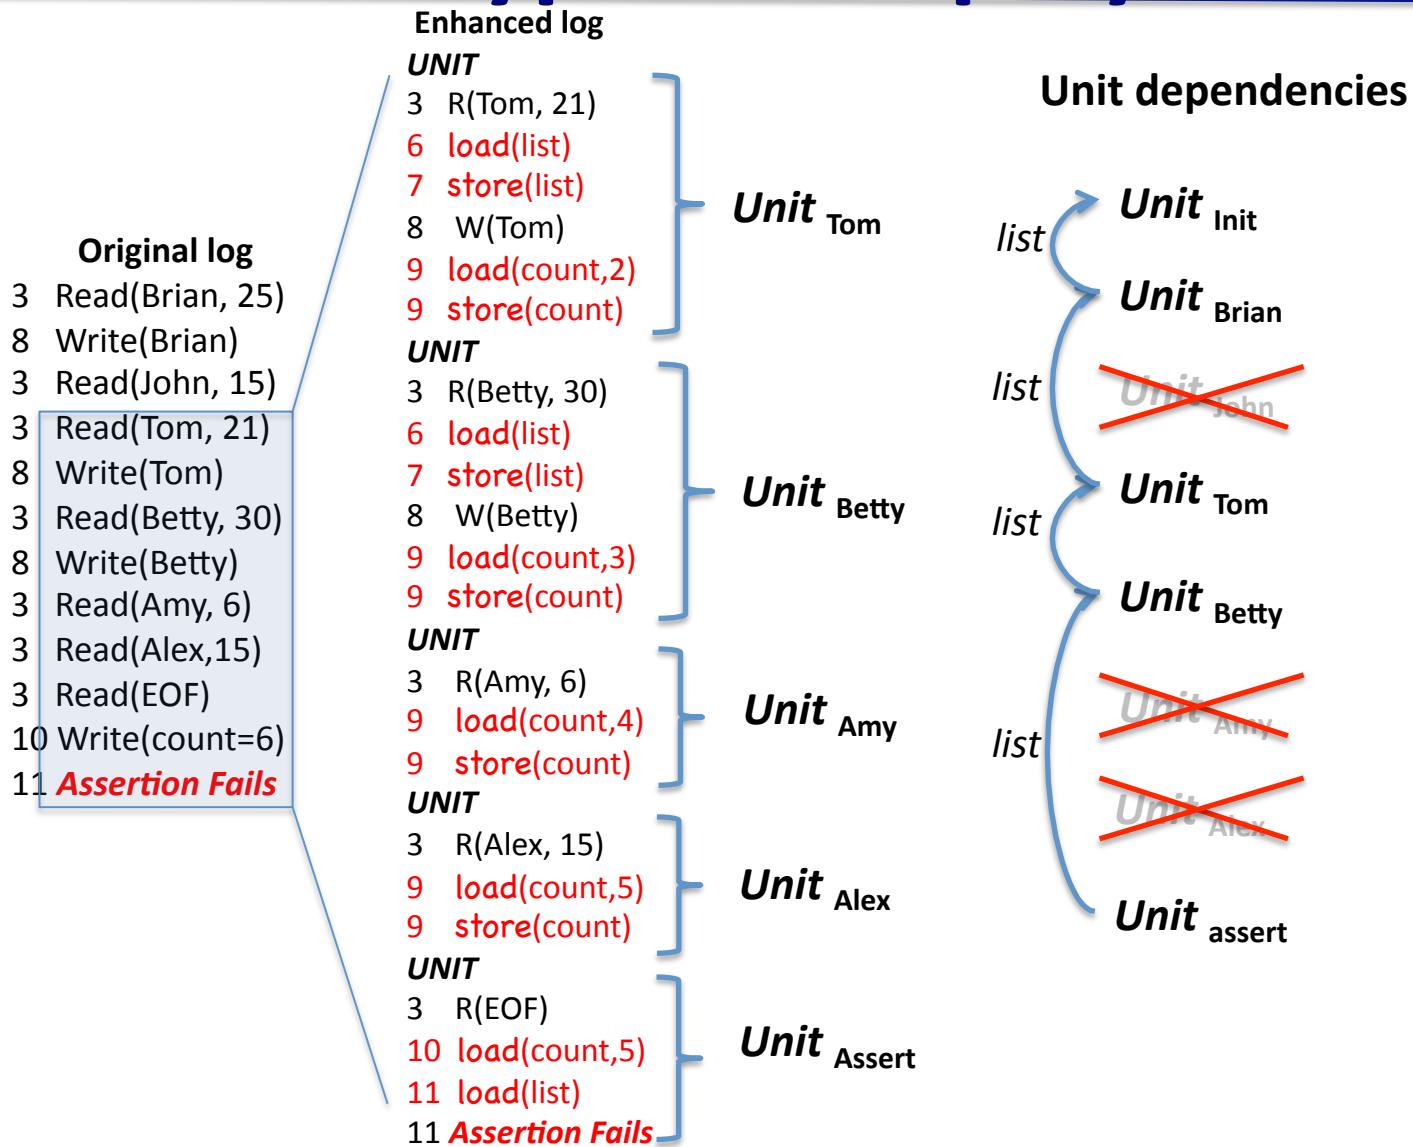


Weave two types of replay





Weave two types of replay

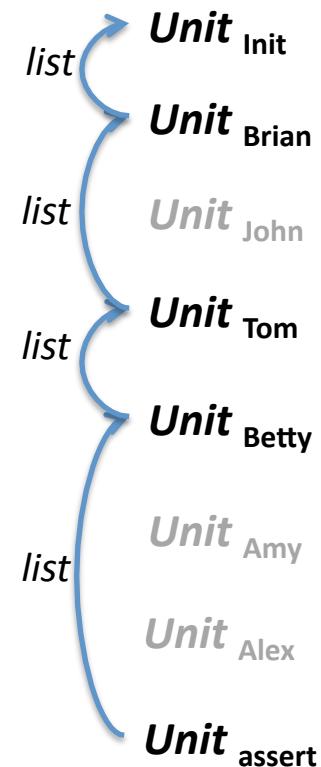




Basic replay

```
1 int count=0;
2 Node *list = null;
3 [UNIT] while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11 }
12 print("count="+count);
13 Assert(list->next==null);
```

Unit dependencies



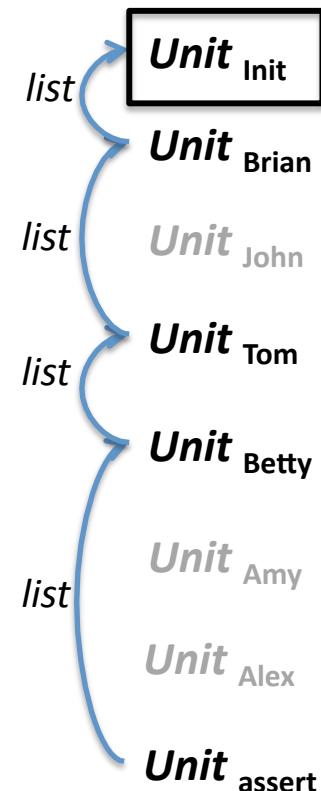


Basic replay

```
1 int count=0;
2 Node *list = null;
3 [UNIT] while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11 }
12 print("count="+count);
13 Assert(list->next==null);
```

list
↓
NULL

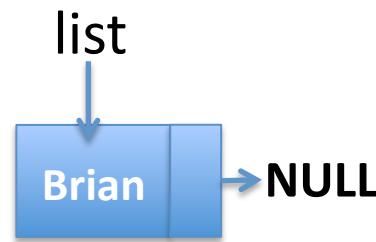
Unit dependencies



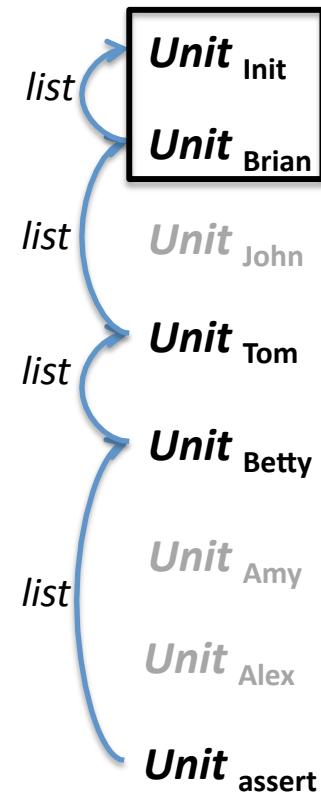


Basic replay

```
1 int count=0;
2 Node *list = null;
3 [UNIT] while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11 }
12 print("count="+count);
13 Assert(list->next==null);
```



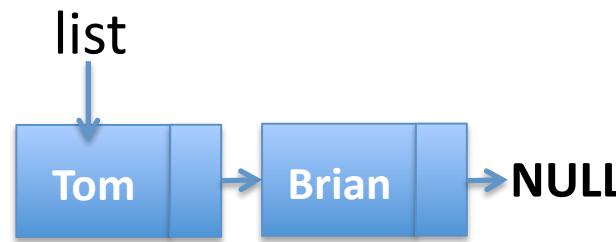
Unit dependencies



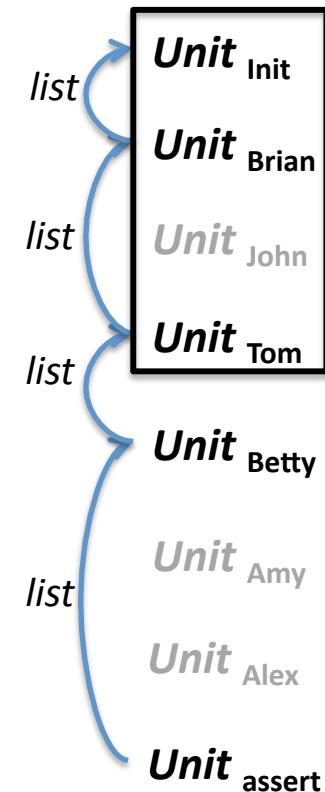


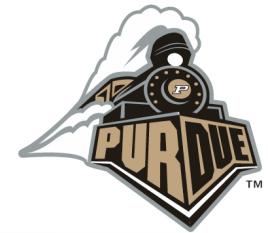
Basic replay

```
1 int count=0;
2 Node *list = null;
3 [UNIT] while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11 }
12 print("count="+count);
13 Assert(list->next==null);
```



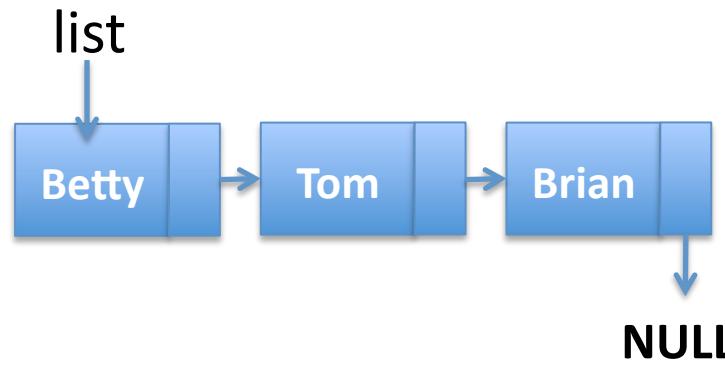
Unit dependencies



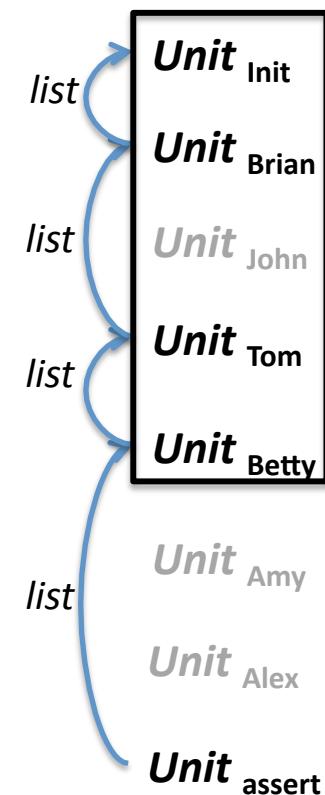


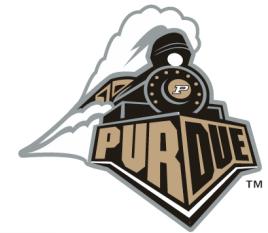
Basic replay

```
1 int count=0;
2 Node *list = null;
3 [UNIT] while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11 }
12 print("count="+count);
13 Assert(list->next==null);
```



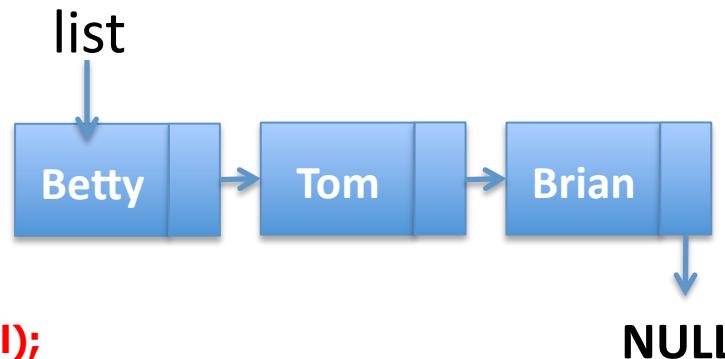
Unit dependencies



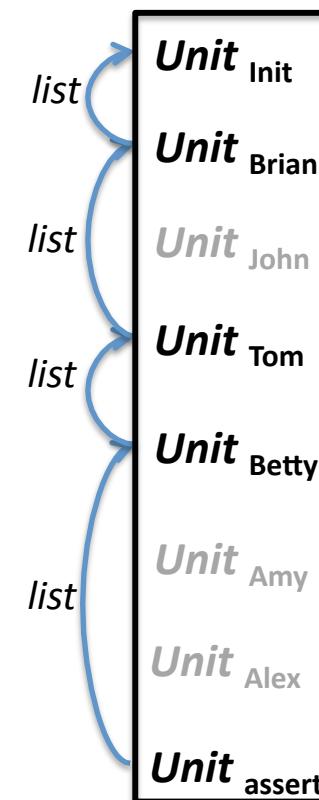


Basic replay

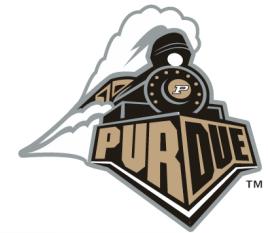
```
1 int count=0;
2 Node *list = null;
3 [UNIT] while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11 }
12 print("count="+count);
13 Assert(list->next==null);
/* Assert the list has one element */
```



Unit dependencies

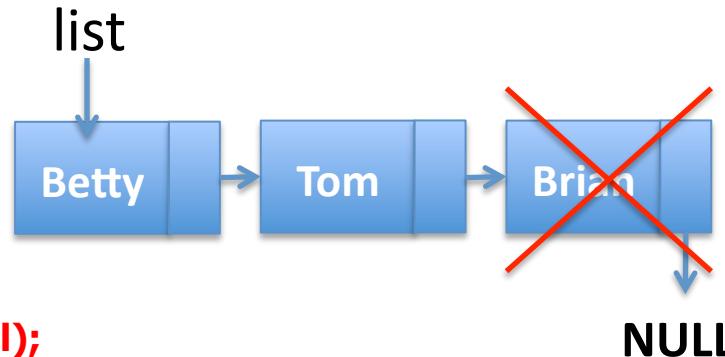


Assertion fail reproduced!

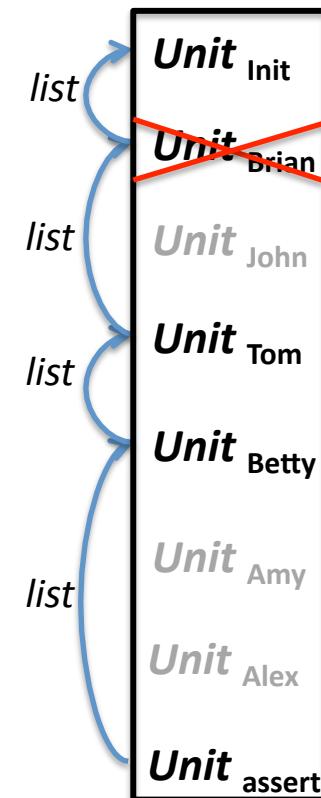


Basic replay

```
1 int count=0;
2 Node *list = null;
3 [UNIT] while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11 }
12 print("count="+count);
13 Assert(list->next==null);
/* Assert the list has one element */
```



Unit dependencies



Assertion fail reproduced!



Aggressive Reduction

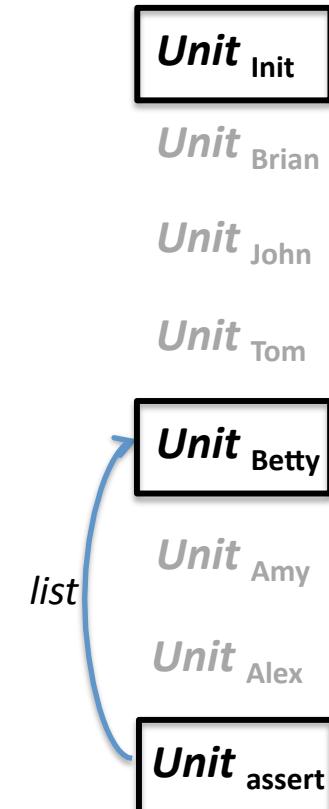
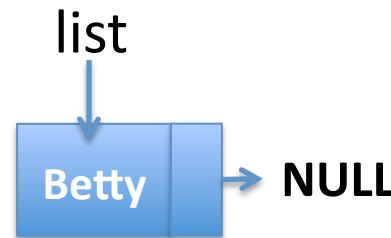
- Compatibility
 - Reconstruct part of execution states, but still can insure replayability
 - We use a BFS algorithm to find a compatible reduction



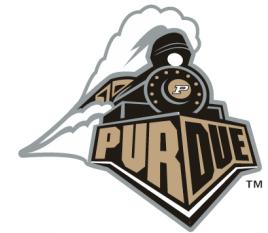
Aggressive Reduction

Unit dependencies

```
1 int count=0;
2 Node *list = null;
3 [UNIT] while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11 }
12 print("count="+count);
13 Assert(list->next==null);
/* Assert the list has one element */
```

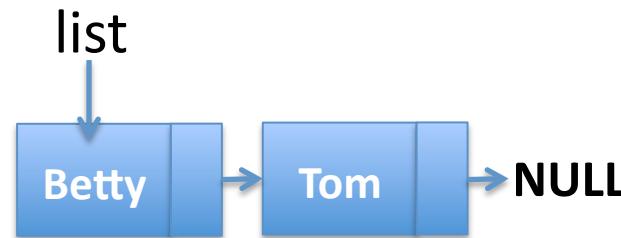


Assertion fail DOES NOT reproduced!

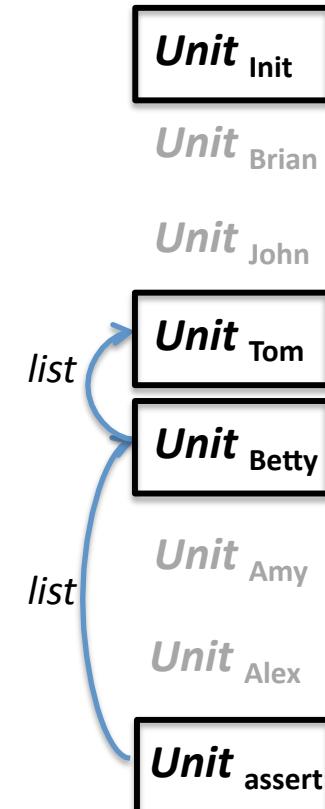


Aggressive Reduction

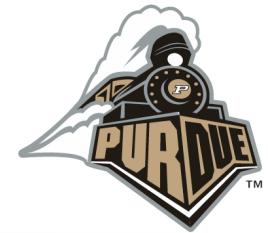
```
1 int count=0;
2 Node *list = null;
3 [UNIT] while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11 }
12 print("count="+count);
13 Assert(list->next==null);
14 /* Assert the list has one element */
```



Unit dependencies



Assertion fail reproduced!

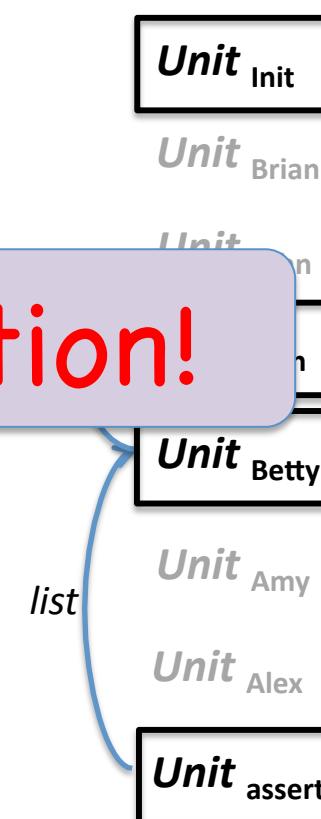


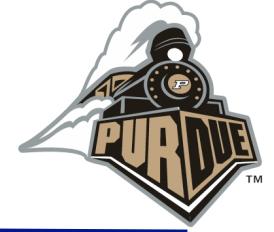
Aggressive Reduction

Unit dependencies

```
1 int count=0;
2 Node *list = null;
3 [UNIT] while(read(&name, &age)) {
4     if(age > 18) {
5         n = new Node(name, age);
6         n->next = list;
7         list = n;
8         print(name);
9     }
10    count++;
11 }
12 print("count=" + count);
13 Assert(list->next==null);
14 /* Assert the list has one element */
```

Compatible Reduction!





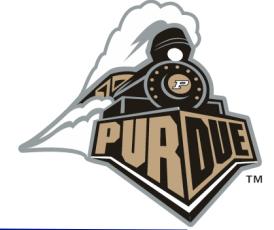
Optimization

UNIT

store(list)

UNIT

load(list)



Optimization

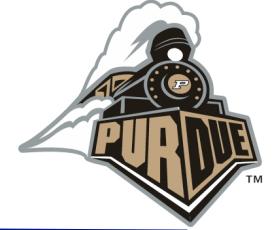
UNIT

store(list)
load(list)

Does not make any
dependence across units!

UNIT

load(list)



Optimization

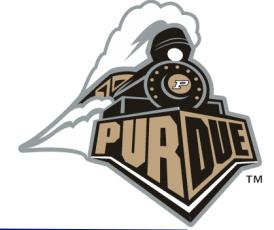
UNIT

store(list)
~~load(list)~~

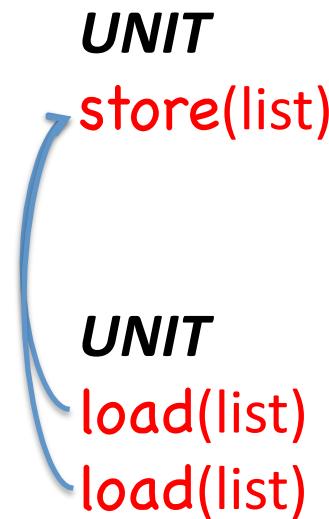
Does not make any
dependence across units!

UNIT

load(list)



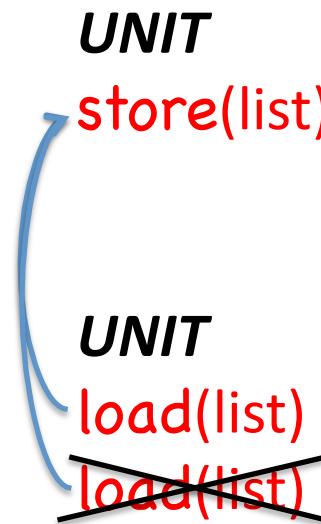
Optimization



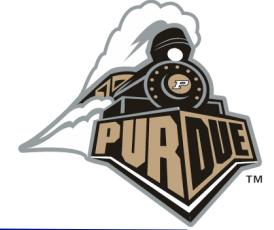
Redundant dependence!



Optimization



Redundant dependence!



Evaluation

- Our system makes use of
 - LLVM – static analysis and instrumentation
 - Jockey – Logging system calls and extra information
 - Pin – Replay component
 - Reduction is implemented in C++
- Evaluate our technique with 11 real bugs from 8 applications



Evaluation – Bug description

Applications	LOC	# of threads	Bug description
Apache-2.0.48	157K	16	#1 : Atomicity violation #2 : Unprotected buffer #3 : Cache size problem
BerkeleyDB-4.7.25	172K	5	#1 : Failure in leader election #2 : Panic caused by out-dated messages
Squid-2.3.4	62K	1	Buffer overflow
MC-4.5.55	106K	1	Buffer overflow
W3M-0.5.2	51K	1	Out of memory
VIM-7.0	230K	1	Hangs (100% CPU usage)
DC-1.3	9.5K	1	Segmentation fault
YAF-C-1.1.1	41K	1	Segmentation fault



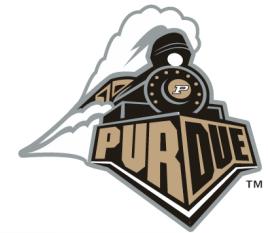
Evaluation – Bug description

Applications	LOC	# of threads	Bug description
Apache-2.0.48	157K	16	#1 : Atomicity violation #2 : Unprotected buffer #3 : Cache size problem
BerkeleyDB-4.7.25	172K	5	#1 : Failure in leader election #2 : Panic caused by out-dated messages
Squid-2.3.4	62K	1	Buffer overflow
MC-4.5.55	106K	1	Buffer overflow
W3M-0.5.2	51K	1	Out of memory
VIM-7.0	230K	1	Hangs (100% CPU usage)
DC-1.3	9.5K	1	Segmentation fault
YAF-C-1.1.1	41K	1	Segmentation fault



Evaluation - Instrumentation

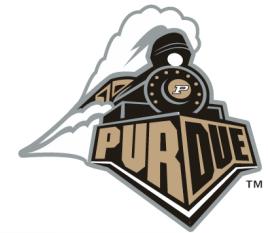
Applications	Annotated loops	Data structure	# of Instrumentation	# of Instrumentation after optimization
Apache-2.0.48	2	32	2,333	1,775
BerkeleyDB-4.7.25	3	47	24,338	15,319
Squid-2.3.4	1	21	2,877	2,232
MC-4.5.55	1	18	1,967	1,460
W3M-0.5.2	1	12	5,506	3,609
VIM-7.0	1	25	12,040	8,375
DC-1.3	1	2	641	521
YAFC-1.1.1	1	8	1,785	1,337
Average	1.37	20.6	6435	4328



Evaluation – Logging overhead

Applications	Time Overhead			Space Overhead (MB)	
	Time (sec)	Jockey overhead	Our overhead	Jockey log	Our log
Apache-2.0.48	210.0	3.46%	0.17%	8.79	47.24
BerkeleyDB-4.7.25	48.3	2.94%	4.20%	6.3	3.43
Squid-2.3.4	79.04	2.53%	1.1%	22.18	30.35
MC-4.5.55	58.1	3.25%	2.26%	12.9	4.34
W3M-0.5.2	N/A	N/A	N/A	4.42	5.09
VIM-7.0	55.26	0.94%	3.93%	4.14	23.27
DC-1.3	35.63	7.11%	4.8%	12.45	5.04
YAFC-1.1.1	63.03	3.48%	1.82%	49.11	1.78

- Overhead by our technique is **2.6% average**.
- Total overhead including Jockey is **5.04 % average**.



Evaluation – Logging overhead

Applications	Time Overhead			Space Overhead (MB)	
	Time (sec)	Jockey overhead	Our overhead	Jockey log	Our log
Apache-2.0.48	210.0	3.46%	0.17%	8.79	47.24
BerkeleyDB-4.7.25	48.3	2.94%	4.20%	6.3	3.43
Squid-2.3.4	79.04	2.53%	1.1%	22.18	30.35
MC-4.5.55	58.1	3.25%	2.26%	12.9	4.34
W3M-0.5.2	N/A	N/A	N/A	4.42	5.09
VIM-7.0	55.26	0.94%	3.93%	4.14	23.27
DC-1.3	35.63	7.11%	4.8%	12.45	5.04
YAFC-1.1.1	63.03	3.48%	1.82%	49.11	1.78

- Overhead by our technique is **2.6% average**.
- Total overhead including Jockey is **5.04 % average**.



Evaluation – Reduction result

Bugs	Log before reduction		Reduced log with search	
	# of units	Log Size(MB)	# of units	Our log(MB)
Apache	15,005	55.84	20	0.86
DB	2,033K	170.6	2	0.06
Squid	8,000	52.6	91	0.69
MC	5,348	17.75	5	1.62
W3M	961	5.65	2	0.051
VIM	2,544	27.41	12	0.314
DC	6,580	17.49	14	1.59
Y AFC	196	113.06	2	0.111
Average	259K	57.55	18.5	0.66



Evaluation – Reduction result

Bugs	Log before reduction		Reduced log with search	
	# of units	Log Size(MB)	# of units	Our log(MB)
Apache	15,005	55.84	20	0.86
DB	2,033K	170.6	2	0.06
Squid	8,000	52.6	91	0.69
MC	5,348	17.75	5	1.62
W3M	961	5.65	2	0.051
VIM	2,544	27.41	12	0.314
DC	6,580	17.49	14	1.59
Y AFC	196	113.06	2	0.111
Average	259K	57.55	18.5	0.66

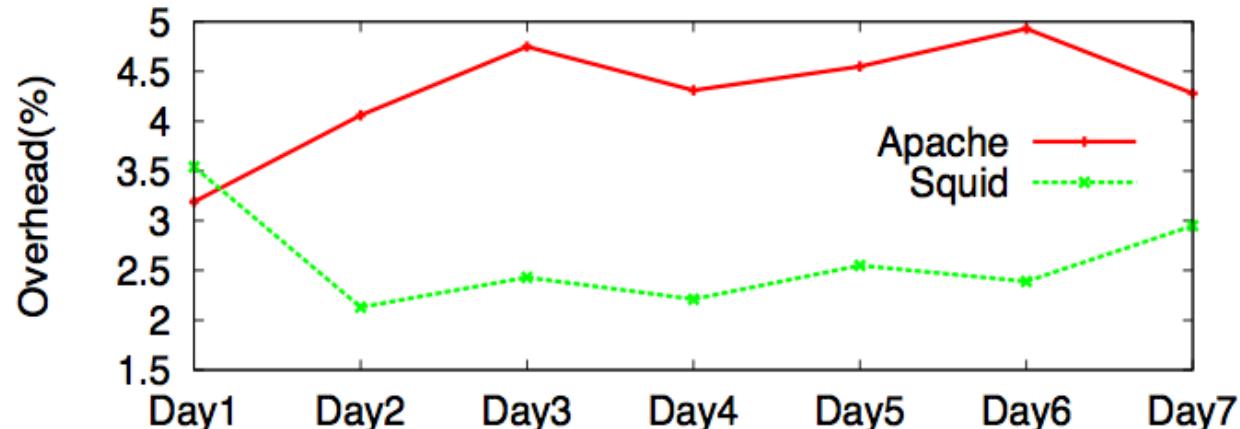


Evaluation – Reduction result

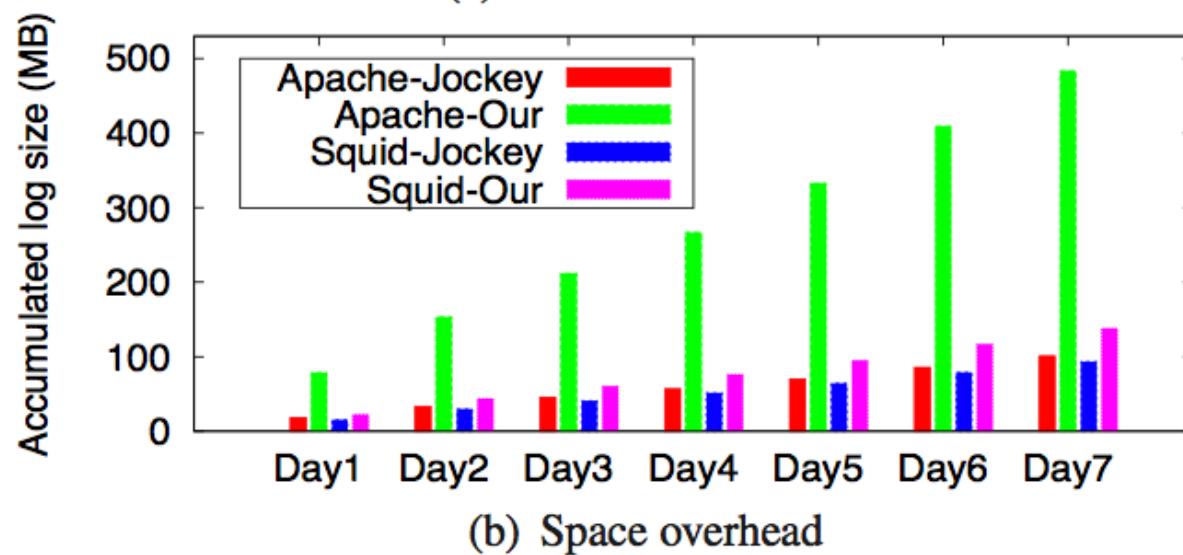
Bugs	Log before reduction		Reduced log with search	
	# of units	Log Size(MB)	# of units	Our log(MB)
Apache	15,005	55.84	20	0.86
DB	2,033K	170.6	2	0.06
Squid	8,000	52.6	91	0.69
MC	5,348	17.75	5	1.62
W3M	961	5.65	2	0.051
VIM	2,544	27.41	12	0.314
DC	6,580	17.49	14	1.59
Y AFC	196	113.06	2	0.111
Average	259K	57.55	18.5	0.66



Evaluation – with real requests



(a) Runtime overhead

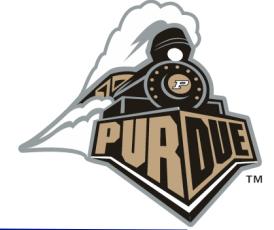


(b) Space overhead



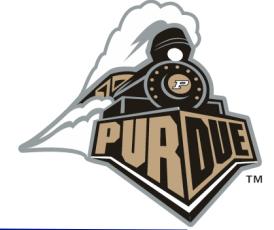
Related Works

- Execution Fast Forwarding
[S.Tallam et al. '07], [X.Zhang et al. '06]
Expensive off-line analysis between events
- Checkpointing
[K.M.Chandy et al. '07], [G. Xu et al. '07],
[G.Broneevetsky et al. '05]
Our technique allows fine-grained reduction



Related Works

- Language based replay system
 - [M.Wu et al. '10], [R.Xue et al. '09]
 - [S.Joshi. '07]
- Software based replay system
 - [G.W.Dunlap et al. '08], [A.Ayers et al. '05]
 - [Z.Guo et al. '08], [S.T.King et al. '05]
- Debugging concurrent programs
 - [G.Altekar et al. '09], [P.Joshi et al. '09],
 - [M.Musuvathi et al. '07], [S.Park et al. '09]



Conclusion

- Compiler based technique that generates a reducible replay log
 - Divides execution into UNITS
 - Instrument programs to collect minimal additional information
 - Reduction can be achieved through analyzing just the log
- Average additional cost by our technique is 2.6%
- Reduce executions with up to 2,033K UNITS to less than 91 Units

Q & A